

Dirac integration with a general purpose bookkeeping DB: a complete general suite for distributed resources exploitation

F Bianchi^{a,b}, M Chrzaszcz^{c,d}, V Ciaschini^e, M Corvo^f, C De Santis^{g,h},
D Del Preteⁱ, A Di Simone^{g,h}, G Donvito^l, A Fella^{m,n}, P Franchini^e, F
Giacomini^e, A Gianelle^f, R Grzymkowski^d, S Longo^f, S Luitz^o, E
Luppi^{p,q}, M Manzali^{p,q}, M Rama^f, G Russoⁱ, S Pardiⁱ, L Perez
Perez^m, B Santeramo^{l,s}, R Stroili^f, L Tomassetti^{n,q}, M Zdybal^d

^a University of Torino, Turin, Italy

^b INFN - Sezione di Torino, Turin, Italy

^c Physik-Institut, Universitat Zurich, Zurich, Switzerland

^d Henryk Niewodniczanski Institute of Nuclear Physics Polish Academy of Sciences, Krakow, Poland

^e INFN - CNAF, Bologna, Italy

^f INFN - Sezione di Padova, Padua, Italy

^g INFN - Sezione di Roma Tor Vergata, Rome, Italy

^h Department of Physics, University of Rome Tor Vergata, Rome, Italy

ⁱ INFN - Sezione di Napoli, Naples, Italy

^l INFN - Sezione di Bari, Bari, Italy

^m INFN - Sezione di Pisa, Pisa, Italy

ⁿ Department of Mathematics and Computer Science, University of Ferrara, Ferrara, Italy

^o SLAC, USA

^p Department of Physics, University of Ferrara, Ferrara, Italy

^q INFN - Sezione di Ferrara, Ferrara, Italy

^r INFN LNF Frascati, Italy

^s Department of Physics, University and Polytechnic of Bari, Bari, Italy

E-mail: bruno.santeramo@ba.infn.it

Abstract. In the context of High Energy Physics computing field the R&D studies aimed to the definition of the data and workload models have been carried on and completed by the SuperB community beyond the experiment life itself. The work resulted of great interest for a generic mid- and small size VO to fulfill Grid exploiting requirements involving CPU-intensive tasks.

We present the R&D line achievements in the design, developments and test of a distributed resource exploitation suite based on DIRAC. The main components of such a suite are the information system, the job wrapper and the new generation DIRAC framework. The DB schema and the SQL logic have been designed to be able to be adaptive with respect to the VO requirements in terms of physics application, job environment and bookkeeping parameters. A deep and flexible integration with DIRAC features has been obtained using SQLAlchemy technology allowing mapping and interaction with the information system. A new DIRAC extension has been developed to include this functionality along with a new set of DIRAC portal interfaces aimed to the job, distributed resources, and metadata management. The results of the first functionality and efficiency tests will be reported.

1. Introduction

In HEP as well in other fields, the necessity to manage and analyze huge amount of data or simulate large amount of events is a typical issue. Today several solutions available for this purpose exist, but lot of them are developed for particular needs of a specific customer. During SuperB R&D activities, a solution useful for SuperB was deployed which can be easily adopted by a generic small and mid-size VO.

2. Suite description

In order to develop a simple, standard and long term solution, suite components (adopted or developed) should be flexible enough to be adapted in order to fulfill needs of a generic VO. Suite components include DIRAC[1], the Information System and the Job Wrapper (see figure 1).

- DIRAC is a well known and widely adopted framework to manage Grid resources, job submission, workflow definition, user authentication, authorization and accounting.
- The Information System holds metadata related to simulations and information about dataset structures and data placement on Grid resources. Information System is a PostgreSQL[2] database. DIRAC interacts with Information System via SQLAlchemy[3].
- The Job Wrapper, executed in bundle with jobs, updates Information System about simulations status and data placement using a REST interface. Job Wrapper is a python script.

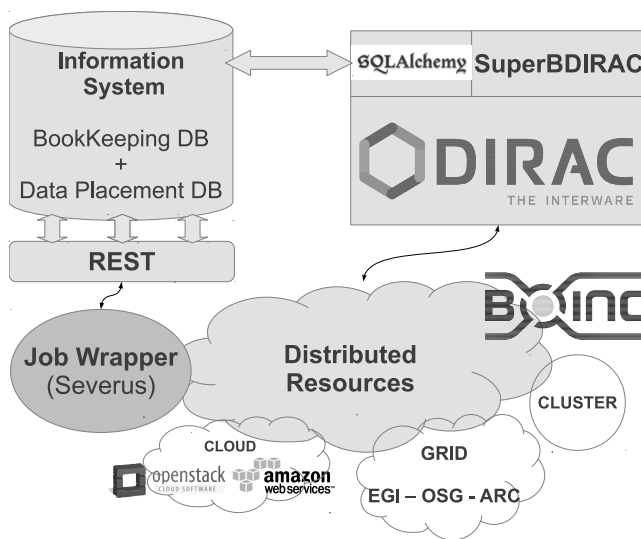


Figure 1. Project bird's eye view.

3. The Dirac extension

DIRAC can be extended to add specific functionalities required by the user. For example in SuperB[5] needs was an interface with a PostgreSQL BookKeeping database and a webportal able to display monitoring data from this database where required. This particular DIRAC extension was named SuperB DIRAC. In SuperB DIRAC a new service, named SBKService, was added to offer an SQLAlchemy layer able to connect DIRAC to a generic SQL DBMS. SBKService maps the bookkeeping database using Object Relational Mapping (see section 5). Webportal extension is designed to create and manage simulations, monitoring related jobs and sites. Integration of new functionalities in webportal permits to use only one interface to manage and monitor the entire stack of simulation related tasks.

4. Bookkeeping DB

MonteCarlo events production needs a method to identify data files and Storage Elements that holds them. A BookKeeping database, named SBK (Super*B* BookKeeping), have been developed to manage metadata associated with data files. The same database stores information about simulations, executed jobs and output data, site availability in terms of installed and supported software. SBK is used also to schedule jobs submission in order to complete simulations. Database is modeled to fulfill general requirements of a typical simulation production.

Entities in SBK are Session, Production and Request. Session defines a simulation (eg. FastSim and FullSim): parameters and software for simulation are defined by VO managers. Production is a Session subset that produce all the needed to simulate a particular scenario (eg. background in detector). Request is a Production subset. The required number of events to complete a Request is defined during its creation. Request completion is monitored via SBK, allowing job re-submission in order to complete it.

SBK design uses the relational model, and the current implementation relies on PostgreSQL (version 9.1) RDBMS which is SQL compliant and, exploiting its hstore datatype, allows to solve some major architectural issues concerning the dataset management of physical parameters. Hstore fields play a major role in the database architecture because storing sets of key/value pairs within a single PostgreSQL value is useful in various scenarios, such as rows with many attributes that are rarely examined, or semistructured data. Beside hstore, some other powerful PostgreSQL features have been exploited: its procedural language (PL/pgSQL) and schemas for a better management of user privileges. An extensive use of views and trigger procedures has been done too.

During the development phase, the SBK database has been continuously analyzed in order to guarantee its normal form (NF) 1, 2 and 3 compliance. Stress test results were the system is capable to sustain 10000 DB transactions (1 transaction = 1 connection+8 insert/update) in 100s (900 operations*sec⁻¹).

5. Bookkeeping DB integration

SBK integration in DIRAC is based on SQLAlchemy. SQLAlchemy uses Object Relational Mapper paradigm: database entities are mapped as python objects. SQLAlchemy adoption simplify code writing, reading and documenting, provides an abstraction layer capable to manage in a transparent way a wide variety of database backends, giving freedom to change it without needs to re-write code. A new DIRAC service, named SBKService, integrates SQLAlchemy functionalities. SBKService interacts with other DIRAC components exactly as any other service. SBKService maps SBK structure and expose methods to perform needed database operations.

6. Job wrapper component

Job wrapper named Severus takes care of main operations of the simulation job: copy software and input files to WN, copy output files in SE and register them in LFC, copy log in SE and bookkeeping DB, update job status in bookkeeping DB.

Access type (lcg or direct) from the worker node to the site SE is automatically recognized and implemented using lcg utils. A module for each session takes care of properly setting environment variables according to simulation (aka session). A configuration file customizes its behavior at execution time.

7. Simulation production use case: the SuperB experience

Simulation Production is designed to manage huge MonteCarlo productions. A webportal, named WebUI[7], is available to manage the entire stack of operations related to this use

case: user management, definition of new "Sessions", "Productions" and "Request", job submission and monitoring, sites management for productions. Production manager users can add and delete sites, add and delete CEs and SEs, set a site as enabled/disabled and supported/unsupported for a given session, manage "Sessions", "Productions" and "Request". Shifter users can submit and monitor jobs for a given Request. Job submission in WebUI is performed via Ganga, while configurations files for submission are generated by WebUI itself taking data from SBK. Ganga submit jobs to grid via WMS using standard gLite commands. Job execution on WNs is driven by Severus (see section 6), while job status updates in SBK is performed via REST interface. Stagein and Stageout as well as LFC registration of output files is performed even by Severus.

SuperBDIRAC goal was the porting of the WebUI functionalities in DIRAC to manage jobs and their submission directly from DIRAC and exploit all the functionalities available in DIRAC: ie. Grid computing resources can be used via WMS or direct submission to CREAM CEs, computer clusters can be accessed via ssh connections, Cloud resources are available via VMDIRAC module, even desktop computers can be used via Boinc. User authentication via X509 certificates and authorization VOMS-role based are builtin in DIRAC. SuperBDIRAC enhance DIRAC integrating the generic bookkeeping database SBK via SQLAlchemy. DIRAC webportal is extended in order to provide an interface for Production manager actions as well as shifter tasks.

8. Functionality Test

A functionality test was performed to demonstrate the capability of SuperBDIRAC to integrate the monitoring of a bookkeeping database in DIRAC webportal.

Functionality Test objective is to demonstrate that SuperBDIRAC is capable of substituting WebUI as monitoring tool for job submission and showing data from a bookkeeping database. Test "Q-factor" is the exact correspondence of information as stored in SBK and displayed in SuperBDIRAC, like in WebUI monitoring page. Correct execution of simulation jobs is not important as long as the bookkeeping information is properly managed.

At present time, not all WebUI functionalities are implemented in SuperBDIRAC, in particular "Submission" job for a given "Request". Therefore the following procedure was followed to perform this test. FastSim job submission is created via WebUI interface. Once submission is created, a set of scripts and configuration files are created whose path is shown by WebUI interface. In particular the php submission script, generated by WebUI, is made of an array with all relevant parameters and UI commands needed to submit jobs in grid via standard glite commands. Since WebUI is still linked with SBK, its portal could be used as well as monitoring portal, useful for a check-cross between info displayed in SBK, WebUI and SuperBDIRAC.

The php submission script is parsed by a python script (`mc_production.py`), in particular the `params` array, in order to retrieve all needed parameters to properly submit jobs via DIRAC: production series, session name, min and max runnumber, configuration files location, physical parameters, events to simulate. Once taken all these parameters, `mc_production.py` uses DIRAC API to prepare and submit jobs via DIRAC client.

DIRAC server receive jobs from DIRAC client, than starts the normal workflow for job management in DIRAC: scheduling, pilot submission, payload retrieval and execution, stageout. Since DIRAC server is equipped with SuperBDIRAC, even bookkeeping monitoring is performed by this component.

Every job simulated 3000 events: this value is set to have an execution time quite longer than 10 minutes. Physical parameters are the same of other official productions. 3 main bunch submission of 400 jobs were performed at INFN-T1 to obtain a total of 1200 simulation jobs. Status in SBK were "prepared", "running" and "done". In addition, 2 bunch submission of 10

jobs were performed, again at INFN-T1, to force some failure messages in SBK and catch it even in SuperB DIRAC monitoring. First failure sample was obtained setting a not-existing site as destination for stageout: error was detected during preliminary check, so status in SBK were "prepared" and "failed". In second failure sample, jobs were submitted using a proxy without Role=ProductionManager, so error occurred in stageout phase: status in SBK were "prepared", "running" and "failed". In summary, 1220 jobs were submitted for test.

| Test | Jobs | prepared | | | running | | | done | | | failed | | | success rate |
|-----------|------|----------|-----|-----|---------|-----|-----|------|-----|-----|--------|----|----|--------------|
| | | A | B | C | A | B | C | A | B | C | A | B | C | |
| good-1 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 0 | 0 | 0 | 100% |
| good-2 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 0 | 0 | 0 | 100% |
| good-3 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 400 | 0 | 0 | 0 | 100% |
| failure-1 | 10 | 10 | 10 | 10 | 0 | 0 | 0 | 0 | 0 | 0 | 10 | 10 | 10 | 100% |
| failure-2 | 10 | 10 | 10 | 10 | 10 | 10 | 10 | 0 | 0 | 0 | 10 | 10 | 10 | 100% |

Table 1. WebUI vs DIRAC comparative test results - A) expected B) in SBK C) in Severus

BookKeeping database was properly updated for every job in every test. All status change were promptly displayed as well in SuperB DIRAC as in WebUI, without any appreciable delay between two portals. SQLAlchemy didn't introduced any appreciable delay or information loss, at least in this functionality test. Table 1 reports, for every test, how many status were saved in SBK and displayed in WebUI and SuperB DIRAC. Success rate was established as ratio between status saved in SBK and status displayed in SuperB DIRAC: its value was 100% in all submissions. SuperB DIRAC could be considered good enough to integrate in DIRAC the capability of monitoring jobs metadata from a bookkeeping database.

9. Conclusions

DIRAC is a mature and stable framework to manage all grid-related tasks, easily adoptable by small as well large VOs. The Information System is designed to be adapted for a generic huge simulation production. The Job Wrapper acts as a bridge between simulation jobs and Information System. We propose SuperB DIRAC as a DIRAC extension capable to satisfy the needs of small and mid size VOs in terms of distributed resource exploitation.

References

- [1] <http://diracgrid.org>
- [2] <http://www.postgresql.org/>
- [3] <http://www.sqlalchemy.org/>
- [4] <http://boinc.berkeley.edu/>
- [5] SuperB Technical Design Report, <http://arxiv.org/abs/1306.5655>
- [6] Fielding R T 2000 *Architectural Styles and The Design of Network-based Software Architectures*, PhD Thesis, University of California Irvine
- [7] A.Fella, E.Luppi, L.Tomassetti *A General Purpose Suite for Job Management, Bookkeeping and Grid Submission*. International Journal of Grid Computing & Applications (IJGCA) Vol.2, No.2, June 2011. DOI: 10.5121/ijgca.2011.2202.